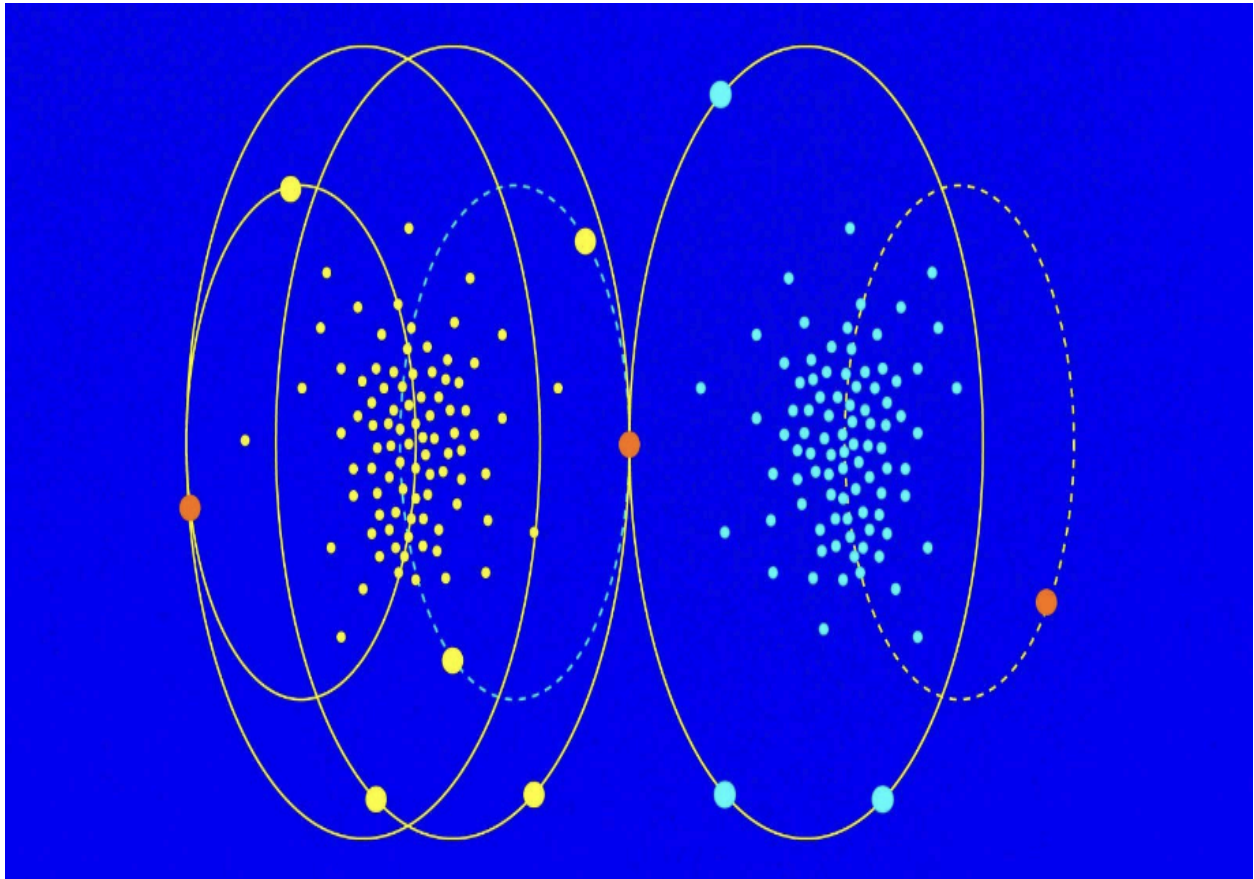


# XCS236 Deep Generative Models

## Course Syllabus



## Welcome

XCS236 Deep Generative Models is a professional course based on graduate-level material from Stanford's on-campus course CS236.

- Learn from Stanford graduate lecture videos (Fall 2023) that have been edited and segmented by topic for easier navigation, reference, and review.
- Complete guided homework assignments implementing content covered in the course lectures.
- Receive support from Stanford-affiliated Course Facilitators.
- Connect to a cohort of peers from diverse locations and professional backgrounds.

## Course Platforms and Tools

**CGOE Learning Management System:** accessed via the [mystanfordconnection](#) site which you used to enroll in this course.

**GitHub:** to distribute code and data for the assignments.

**Slack:** for additional course support and class discussions.

## Important Dates

**January 24 afternoon Pacific Time: Slack and GitHub invites sent** (Accept within 7 days)

→ Did not receive the Slack invite and it's ***not in your spam?***

- <http://XCS236.slack.com/>
- I have a Guest Account
- Log in using your credentials

→ Did not receive the GitHub Invite?

- Email [xcs236-staff@stanford.edu](mailto:xcs236-staff@stanford.edu)
- Let us know if you'd like to receive the invite to a different email

**January 27 noon Pacific all lecture videos become available.**

- Log in to your [mystanfordconnection](#) account
- Click on the link titled "Course Videos and Assignments"

→ Here is a video on how to navigate the course portal.

**January 27 evening Pacific Time Course Facilitator connection emails sent out to learners.**

→ Course Facilitators serve as your primary point of contact for content and assignment-related questions

**February 14 at 5:00pm PST Drop/Transfer Deadline**

## Deadlines and Pacing

**Course Start:** January 27

**Course End:** April 6

Below is a *suggested* pacing guide. Please, note the assignment deadlines (all **11:59 PM Pacific**).

Week	Suggested Videos	Assignments		
		Release Date	Regular Deadline	Late Deadline
1 & 2 Jan 27 - Feb 9	Introduction: Deep Generative Models PyTorch Tutorial Introduction to Generative Models Autoregressive Models Maximum Likelihood Learning Latent Variable Models Variational Autoencoder (VAE)	Assignments 1-3 January 27		
3 & 4 Feb 10 - Feb 23	Normalizing Flow Models Advanced GAN Topics Generative Adversarial Networks (GANs)		Assignment 1 February 16	Assignment 1 February 21
5 & 6 Feb 24 - Mar 9	Energy Based Models (EBMs) Score Based Models	Assignment 4 March 3	Assignment 2 March 2	Assignment 2 March 7
7 & 8 Mar 10 - Mar 23	Score Based Generative Modeling Evaluating Generative Models		Assignment 3 March 16	Assignment 3 March 21
9 & 10 Mar 24 - Apr 6	Score Based Diffusion Models Diffusion Models for Discrete Data		Assignment 4 April 6	Assignment 4 April 11

## Late Deadlines and Penalty Waiver

**Late Deadlines:** All assignment submissions can be turned in up to five days late and are assessed a penalty of -1 point per late day. After five days, the submission link will close, and entries will no longer be accepted.

**Penalty Waiver:** You have the option to remove late penalty points from any one assignment.

**To request a penalty waiver**, please complete the following google form [linked here](#). Please ensure that the email address used for your submission matches the one you used at enrollment. All requests will be applied at the end of the course.

## Assignments and Grading

[Here is a short description of each assignment.](#)

**Coding Questions:** are graded automatically upon upload and can be submitted up until the late deadline. [Click here for an example.](#)

**Written Questions:** will be manually graded by Course Facilitators no later than one week after the assignments 'on-time' deadline. [Click here for an example of the submission process.](#)

Note: All assignments require prior knowledge of machine learning principles.

Learners may submit regrade requests on Gradescope within one week of the assignment grades being posted in cases of clear grading errors, such as miscalculations.

## Certificate Requirements

The course is pass/no-pass, and no letter grades are granted. To pass the course, you must achieve a total score of 70% or higher on the assignments. Upon successful completion, you will receive a digital course certificate.

- There are a total of 200 base points (meaning 140 to achieve 70%).
- Points distribution table will be added soon.

**Note:** Please note that the course certificates will only be sent at the end of the course.

## Course Facilitators and Support

We encourage posting any class/content-related questions in relevant Slack channels. This way, you'll likely get a faster response from either the course staff or your classmates. However, before you post, make sure you're familiar with the important [course policies](#).

Additionally, starting from the first day of the course, you'll be connected to a Course Facilitator (CF), who will be your primary point of contact for content and assignment inquiries. They'll each lead a smaller group of learners, providing personalized support. Your Course Facilitator will keep you informed about important reminders and their availability for questions, potential online office hours, and 1:1 sessions.

## Note on Code Assignments and Debugging

While the course team is here to support your experience, it is ultimately your responsibility to write, test, and debug your own coding assignments. Before reaching out to a CF or posting your question in Slack, it's expected that you have taken the reasonable step of performing an analysis yourself. CFs may view and provide guidance on your work; however, they will not provide exact answers on what to insert into your assignments. This policy is meant to ensure that you leave the course having mastered the material and enables CFs to focus on questions where their guidance is most impactful.

## Drop/Transfer Policy

We don't want to see you go, but if you decide this is not the right course or time, there are two available options: either drop the course OR transfer to the next iteration of this course/another course within the AI Professional Program. All learners have the option to transfer a maximum of one (1) time per course for their duration in the program. Once this one-time transfer is used for a specific course, no additional transfer requests for that course will be accepted.

To request a drop or transfer, email [ai-drop-transfer@stanford.edu](mailto:ai-drop-transfer@stanford.edu). Please include the course number and request type (drop or transfer) in the subject line to ensure your request is processed correctly.

<b>Up until</b>	<b>January 27</b>	No cost for drop/transfer. If you drop, you will get a full refund.
<b>Up until</b>	<b>February 14</b>	Once the course has begun, there will be a drop/transfer fee of \$200, i.e.: If you request a drop, you will be reimbursed 100% of your tuition minus \$200. If you request a transfer, there will be a \$200 fee in the form of an invoice.

**Important Note:** Beyond the third week of the course, tuition fees are not granted

## Important Policies

### Honor Code

Students will be asked to review and maintain the standard set forth by the [Stanford Honor Code](#) when completing quizzes and assignments in this course. You can review the section labeled Violations of the Honor Code for representative examples relevant to this course.

Students are strongly encouraged to form study groups, discuss, and work on homework problems in groups, and help each other; However, each student must write down the solutions independently and cannot refer to written notes from the joint session. **In other words, you must understand the solution well enough in order to reconstruct it independently.** Further, because we occasionally reuse assignment questions from previous years, you are expected not to copy, refer to, or look at the solutions in preparing your answers. It is an honor code violation to intentionally refer to a previous year's solutions

After completing this course, you are welcome to share your experience and credentials with others; however, **it is considered a violation of the honor code to share assignment solutions including on public platforms such as GitHub.** Faculty in the computer science department have strongly encouraged us to refrain from posting solutions for assignments, thus we ask that you **DO NOT** share the exact code.

### Communication Guidelines

One of the benefits of this course is the opportunity to network with other course participants, to create study groups and to engage with course staff. We encourage this kind of interaction and want to make sure that it is a positive experience for everyone. It is imperative that no course participant is made to feel uncomfortable or their ability to learn or otherwise benefit from the course is impeded by the actions of another participant. Please use good judgment. Keep all interactions professional and focused on coursework or career networking. Maintain respect in all your interactions by using polite language, avoid using offensive language and be considerate of your colleagues' preferences regarding direct messaging. Please respect and uphold the rights and dignity of others regardless of race, color, national or ethnic origin, sex, age, disability, religion, sexual orientation, gender identity, or socio-economic status. Our team is always available either in Slack or via email, so please feel free to reach out to us if you have any questions or concerns.

You can review Stanford Online's [terms of service](#) here, including rules for online conduct.